

AMMM Course Project

*Optimal Priority Assignment in
Cooperative GPU Scheduling*

**Ezgi Sena Karabacak
Davide Lamagna**

Problem Overview

- A cooperative of N users shares a GPU
- Users bid to gain priority over others on overlapping days
- **Objective:** Maximize collected bids while avoiding cycles (deadlocks)
- **Key constraint:** Resulting priority graph must be **acyclic**

Formal Problem Definition

Inputs:

- Number of members N
- Bid matrix m_{ij}
- Bids are only meaningful if members i and j have overlapping requests.

Output:

- A directed acyclic graph indicating priority relations

Goal:

- Maximize total bid value without forming cycles

ILP Model

Decision variables:

- $x_{ij} \in \{0, 1\} \rightarrow 1$ if member i has priority over j .
- $u_i \in \{1, 2, \dots, N\} \rightarrow$ topological rank of node i

Objective:

$$\text{Maximize } \sum_{i \neq j} x_{ij} \cdot m_{ij}$$

Constraints:

- No self-priority: $x_{ii} = 0$ for all i
- Acyclicity via topological order:
If $x_{ij} = 1$, then $u_i < u_j$
(Enforced with: $u_i + 1 \leq u_j + (1 - x_{ij}) * N$)

Heuristic Algorithms

Why heuristics?

- CPLEX gives optimal results but becomes slow or times out for large instances ($N > 45$).
- We need **faster**, **scalable** alternatives that still give good-quality solutions.

Implemented Heuristics:

- Greedy Constructive Heuristic
- Greedy + Local Search
- GRASP
- GRASP + Local Search

Goal: Find near-optimal solutions much faster than CPLEX, especially for large problem sizes.

Greedy Constructive Heuristic

- Sort all possible arcs (i, j) by bid value m_{ij} in descending order.
- Start with an empty graph.
- Add each arc one by one if it does not create a cycle.
- Continue until all arcs have been checked.

Advantages:

- Very fast.
- Produces a valid (acyclic) solution.
- Quality depends heavily on early decisions.

Limitation:

- Can miss better configurations due to early greedy choices.

Greedy Pseudocode

Algorithm 1: Greedy Constructive Heuristic

Input: Bid matrix m_{ij}

Output: Acyclic orientation A

Initialize $A \leftarrow \emptyset$;

Add all nodes $i \in \{1, \dots, N\}$ to A ;

Sort all arcs (i, j) with $i \neq j$ by m_{ij} in descending order;

for *each* (i, j) *in sorted list* **do**

 Temporarily add edge (i, j) to A ;

if A *remains acyclic* **then**

 Permanently keep edge (i, j) ;

else

 Remove edge (i, j) from A ;

return A

Local Search Strategy

- Improves the greedy solution by modifying the topological order.
- Swaps each node with nearby nodes (+1, +2, +3) to find better bid values.
- Only keeps swaps that maintain acyclicity and increase total value.
- Uses **Best Improvement**: chooses the best swap per iteration.
- Stops when no more improving moves are found or timeout is reached.

Local Search Pseudocode

Algorithm 2: Local Search Algorithm

Input: Initial solution A , bid matrix m_{ij}

Output: Improved solution A'

Initialize $order \leftarrow$ topological sort of A ;

Compute initial fitness $f \leftarrow \sum_{(i,j) \in A} m_{ij}$;

repeat

for each $i \in \{1, \dots, N - 1\}$ **do**

for $\delta \in \{1, 2, 3\}$ **do**

$j \leftarrow i + \delta$;

if $j < N$ **then**

 Swap $order_i$ and $order_j$;

 Construct new graph A' from new order;

 Compute $f' \leftarrow \sum_{(u,v) \in A'} m_{uv}$;

if A' is acyclic and $f' > f$ **then**

 Accept new order and update fitness $f \leftarrow f'$;

until no improvement or timeout;

return A'

GRASP Heuristic

- Builds solutions iteratively using:
 - A Restricted Candidate List (RCL) based on bid values.
 - Random selection from the RCL (controlled by parameter α).
- After construction, applies Local Search for refinement.

Key Benefit:

- Combines exploration (randomness) and exploitation (local search).
- Avoids greedy traps and finds better-quality solutions.

Tuned Parameter:

$\alpha = 0.1$ gave the best performance in our experiments.

GRASP Pseudocode

Algorithm 3: GRASP with Local Search

Input: Bid matrix m_{ij} , parameter α , time budget

Output: Best acyclic orientation A^*

Initialize $A^* \leftarrow \emptyset$, $f^* \leftarrow 0$;

repeat

 Build RCL from current candidates based on α ;

 Randomly select $(i, j) \in \text{RCL}$;

 Add (i, j) to A if no cycle is formed;

 Repeat until no more candidates;

if *local search is enabled* **then**

 └ Apply Local Search to improve A to get A' with fitness f ;

else

 └ $A' \leftarrow A$

if $f > f^*$ **then**

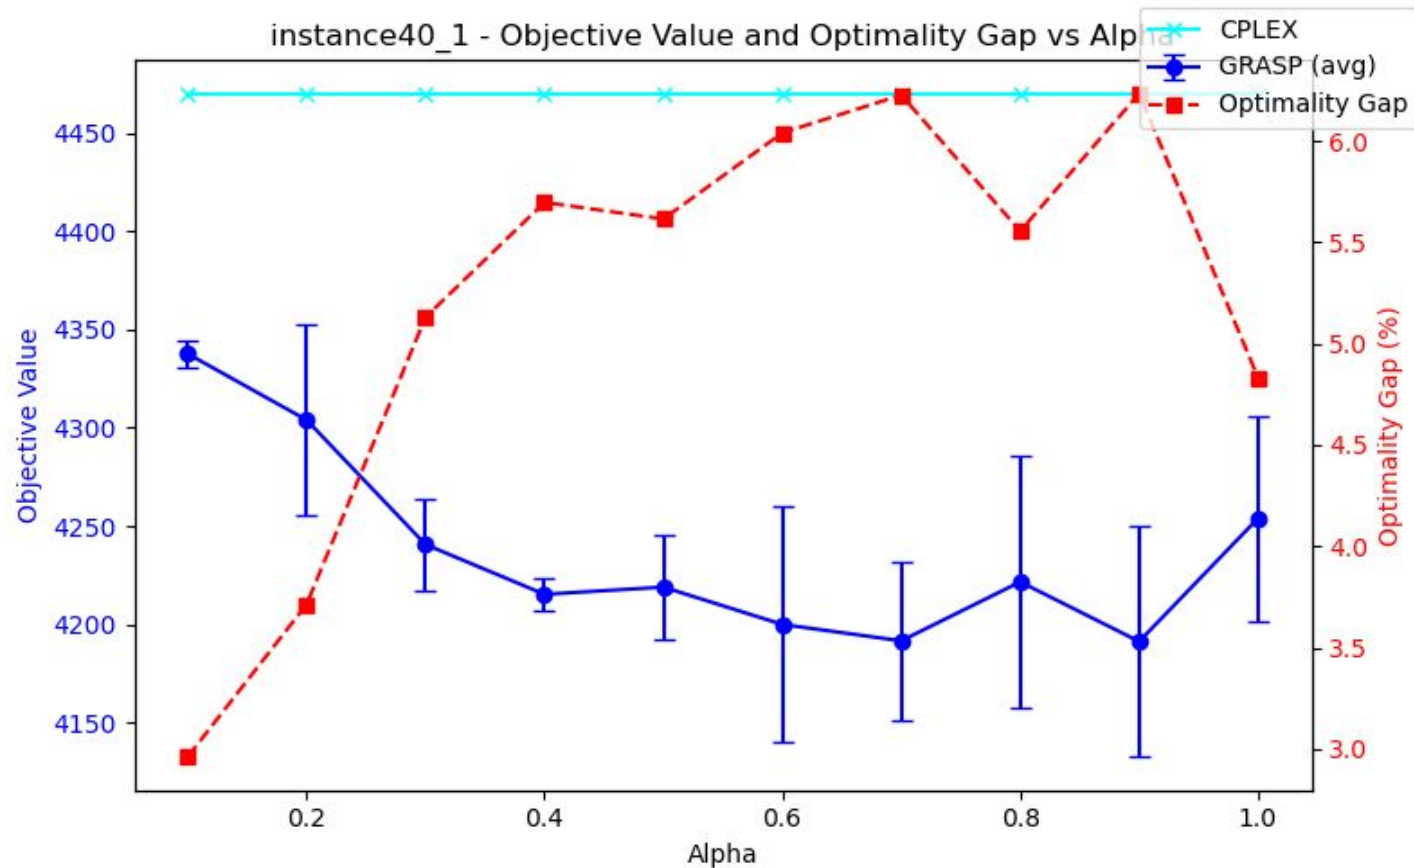
 └ Update $A^* \leftarrow A'$, $f^* \leftarrow f$;

until *stopping criteria met*;

return A^*


Alpha Tuning

N=40



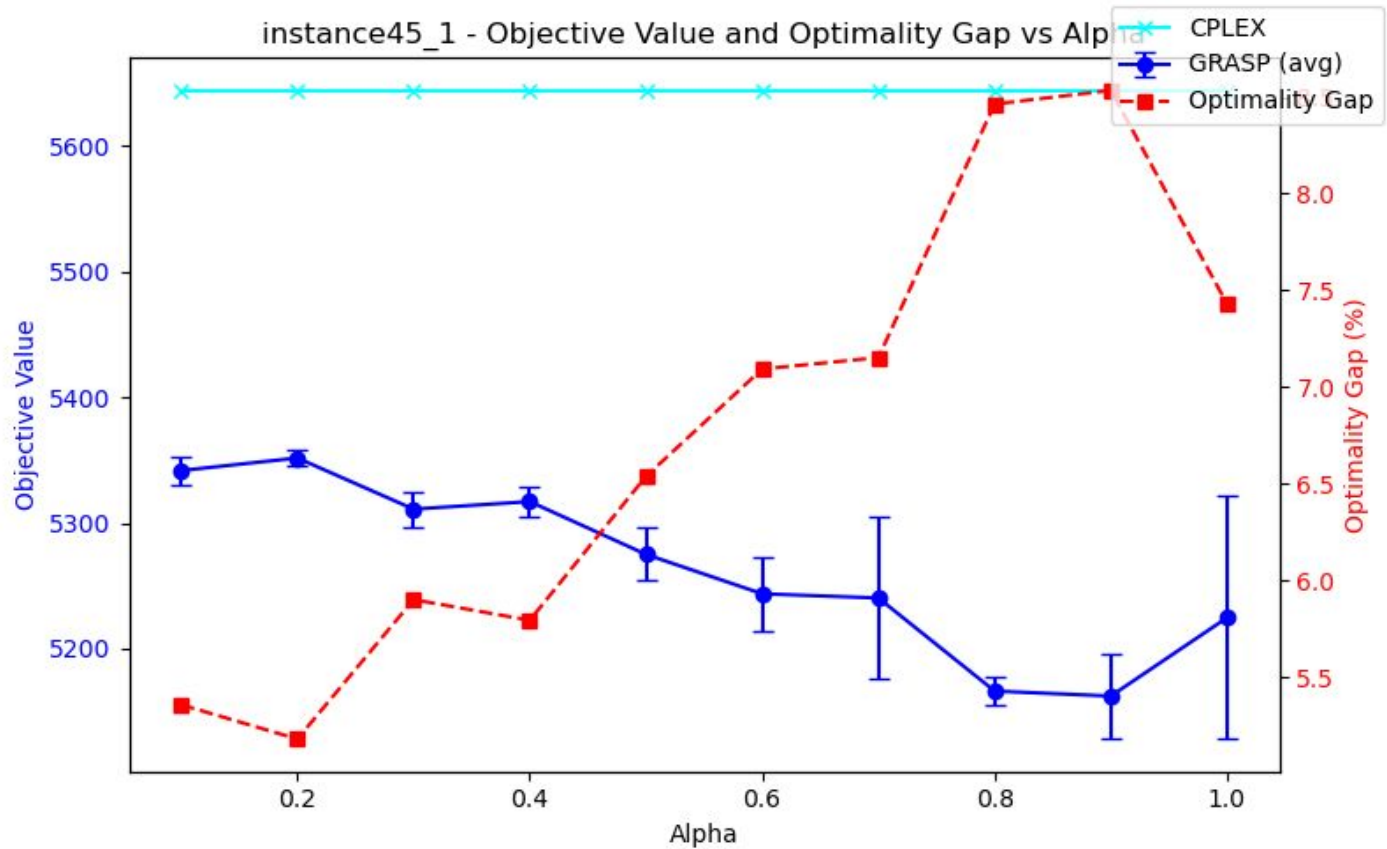
Alpha Tuning

N=43



Alpha Tuning

N=45



Experimental Setup


- Instance sizes tested: **N = 40** to **N = 50**
- Bid values: Random integers in the range [1, 10]
- GRASP tuning:
 - $\alpha \in \{0.1, 0.2, \dots, 1.0\}$
 - Each value tested 3 times
 - Selected $\alpha = 0.1$ based on best average performance
- CPLEX timeout: 60 seconds
- Starts to fail beyond **N = 45**

All results include:

- Objective value
- Elapsed time
- Iterations (for heuristics)


Results

Objective value vs N



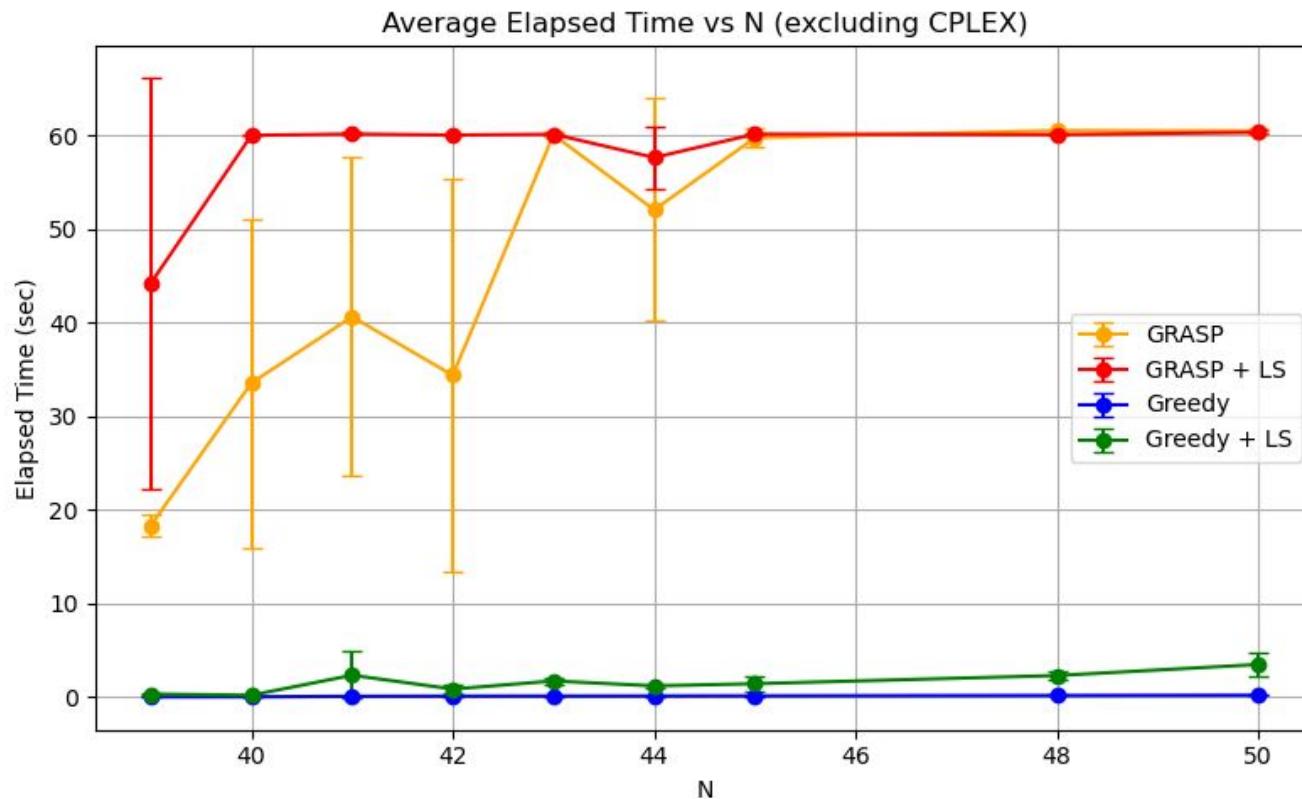
Results

Elapsed Time vs N




Results

Elapsed Time vs N



Results

Optimality Gap vs N



Conclusion

- CPLEX gives the best solutions but becomes impractical for large instances
- GRASP + Local Search offers the best trade-off:
 - Near-optimal solutions (within 4–5% of CPLEX)
 - Much faster and scalable
- Greedy + Local Search is a simple and fast alternative with reasonable quality.
- Heuristic methods are effective for large-scale or time-sensitive problems.

Thank you :)

Any questions?